

fully conv autoencoder

March 24, 2019

```
In [1]: from matplotlib import pyplot as plt
import numpy as np
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Activation, Flatten, MaxPooling1D, UpSampling1D
from keras.layers.convolutional import Conv1D, Conv2D
import medleydb as mdb
from scipy import signal
from scipy.io import wavfile
```

Using TensorFlow backend.

1 Load the mix

```
In [65]: multitricks = mdb.load_multitracks(["Phoenix_ColliersDaughter"])
```

```
In [66]: mix = next(multitricks)
```

```
In [67]: sample_rate, mix_audio = wavfile.read(mix.mix_path)
mix_audio = mix_audio.mean(1)
mix_audio = mix_audio[sample_rate*8:int(-sample_rate*1.5)]
```


```
In [68]: samples_per_period = 256
overlap = int(samples_per_period/2)
```

```
In [69]: freqs, times, s_mix = signal.stft(mix_audio, fs=sample_rate, nfft=samples_per_period,
noverlap=overlap, nperseg=samples_per_period)
```

```
freqs = freqs[:-1]
s_mix = s_mix[:-1,:]
```

```
In [70]: plt.pcolormesh(times, freqs, 20*np.log10(np.abs(s_mix)))
```

```
Out[70]: <matplotlib.collections.QuadMesh at 0x1b498fe6630>
```



2 Load the Flute

```
In [71]: flutes = mdb.get_files_for_instrument("flute", [mix])
```

```
In [72]: flute = next(flutes)
```


```
In [73]: sample_rate, flute_audio = wavfile.read(flute)
         flute_audio = flute_audio.mean(1)
         flute_audio = flute_audio[sample_rate*8:int(-sample_rate*1.5)]
```

```
In [74]: freqs, times, s_flute = signal.stft(flute_audio, fs=sample_rate, nfft=samples_per_peri
                                             noverlap=overlap, nperseg=samples_per_peri
```

```
         freqs = freqs[: -1]
         s_flute = s_flute[: -1, :]
```

```
In [75]: plt.pcolormesh(times, freqs, 20*np.log10(np.abs(s_flute)))
```

```
Out [75]: <matplotlib.collections.QuadMesh at 0x1b487944978>
```




2.1 Create a mask for the flute

In [76]: `mask_flute = s_flute / (s_flute + s_mix + 1e-9)`

In [77]: `plt.pcolormesh(times, freqs, 20*np.log10(np.abs(mask_flute)))`

Out [77]: `<matplotlib.collections.QuadMesh at 0x1b4879a9710>`



3 Prepare the training data

```
In [78]: num_test = int((sample_rate * 10)/samples_per_period)
        num_train = s_mix.shape[1] - num_test
        print(num_train, num_test)
```

60694 1722

```
In [79]: mix_train = s_mix.T[:num_train,:,np.newaxis]
        mix_train = np.concatenate((mix_train.real, mix_train.imag), axis=-1)
        mix_train.shape
```

Out[79]: (60694, 128, 2)

```
In [80]: flute_train = mask_flute.T[:num_train,:,np.newaxis]
        flute_train = np.concatenate((flute_train.real, flute_train.imag), axis=-1)
        flute_train.shape
```

Out[80]: (60694, 128, 2)

```
In [81]: mix_test = s_mix.T[-num_test:,:,np.newaxis]
        mix_test = np.concatenate((mix_test.real, mix_test.imag), axis=-1)
        mix_test.shape
```

Out[81]: (1722, 128, 2)

```
In [82]: flute_test = mask_flute.T[-num_test:,:,np.newaxis]
        flute_test = np.concatenate((flute_test.real, flute_test.imag), axis=-1)
        flute_test.shape
```

Out[82]: (1722, 128, 2)

```
In [83]: model = Sequential()
```

```
model.add(Conv1D(16, 2, padding="same", input_shape=mix_train.shape[1:], activation="relu"))
model.add(MaxPooling1D(2, padding="same", name="MaxPooling_1"))
```

```
model.add(Conv1D(8, 2, padding="same", activation="relu", name="Conv1D_2"))
model.add(MaxPooling1D(2, padding="same", name="MaxPooling_2"))
```

```
model.add(Conv1D(4, 2, padding="same", activation="relu", name="Conv1D_3"))
model.add(MaxPooling1D(2, padding="same", name="MaxPooling_3"))
```

```
model.add(Conv1D(4, 2, padding="same", activation="relu", name="Conv1D_4"))
model.add(UpSampling1D(2, name="UpSampling_1"))
```

```

model.add(Conv1D(8, 2, padding="same", activation="relu", name="Conv1D_5"))
model.add(UpSampling1D(2, name="UpSampling_2"))

model.add(Conv1D(16, 2, padding="same", activation="relu", name="Conv1D_6"))
model.add(UpSampling1D(2, name="UpSampling_3"))

model.add(Conv1D(2, 2, padding="same", name="Conv1D_7"))

model.summary()

```

Layer (type)	Output Shape	Param #
Conv1D_1 (Conv1D)	(None, 128, 16)	80
MaxPooling_1 (MaxPooling1D)	(None, 64, 16)	0
Conv1D_2 (Conv1D)	(None, 64, 8)	264
MaxPooling_2 (MaxPooling1D)	(None, 32, 8)	0
Conv1D_3 (Conv1D)	(None, 32, 4)	68
MaxPooling_3 (MaxPooling1D)	(None, 16, 4)	0
Conv1D_4 (Conv1D)	(None, 16, 4)	36
UpSampling_1 (UpSampling1D)	(None, 32, 4)	0
Conv1D_5 (Conv1D)	(None, 32, 8)	72
UpSampling_2 (UpSampling1D)	(None, 64, 8)	0
Conv1D_6 (Conv1D)	(None, 64, 16)	272
UpSampling_3 (UpSampling1D)	(None, 128, 16)	0
Conv1D_7 (Conv1D)	(None, 128, 2)	66
Total params: 858		
Trainable params: 858		
Non-trainable params: 0		

```
In [84]: model.compile('adam', loss='mean_squared_error', metrics=['accuracy'])
```

```
In [85]: history = model.fit(mix_train, flute_train, batch_size=200, epochs=10)
```

```

Epoch 1/10
60694/60694 [=====] - 63s 1ms/step - loss: 4.3759 - acc: 0.7718
Epoch 2/10
60694/60694 [=====] - 11s 179us/step - loss: 4.1513 - acc: 0.8055
Epoch 3/10
60694/60694 [=====] - 12s 194us/step - loss: 4.1494 - acc: 0.8079
Epoch 4/10
60694/60694 [=====] - 12s 192us/step - loss: 4.1487 - acc: 0.8090
Epoch 5/10
60694/60694 [=====] - 11s 183us/step - loss: 4.1483 - acc: 0.8095
Epoch 6/10
60694/60694 [=====] - 11s 184us/step - loss: 4.1479 - acc: 0.8099
Epoch 7/10
60694/60694 [=====] - 11s 179us/step - loss: 4.1477 - acc: 0.8103
Epoch 8/10
60694/60694 [=====] - 11s 184us/step - loss: 4.1475 - acc: 0.8104
Epoch 9/10
60694/60694 [=====] - 11s 181us/step - loss: 4.1474 - acc: 0.8105
Epoch 10/10
60694/60694 [=====] - 11s 181us/step - loss: 4.1473 - acc: 0.8106

```

```
In [86]: results = model.evaluate(mix_test, flute_test)
```

```

print(model.metrics_names)
print(results)

```

```

1722/1722 [=====] - 1s 359us/step
['loss', 'acc']
[0.5205799524653941, 0.8178308290471599]

```

4 Results

```
In [87]: mask_prediction = model.predict(mix_test)
        mask_prediction.shape
```

```
Out[87]: (1722, 128, 2)
```

```
In [88]: mask_prediction = np.sqrt((mask_prediction**2).sum(-1)).T
```

```
In [89]: fig, (ax1, ax2, ax3) = plt.subplots(1,3,sharey=True,figsize=(12, 4))
```

```

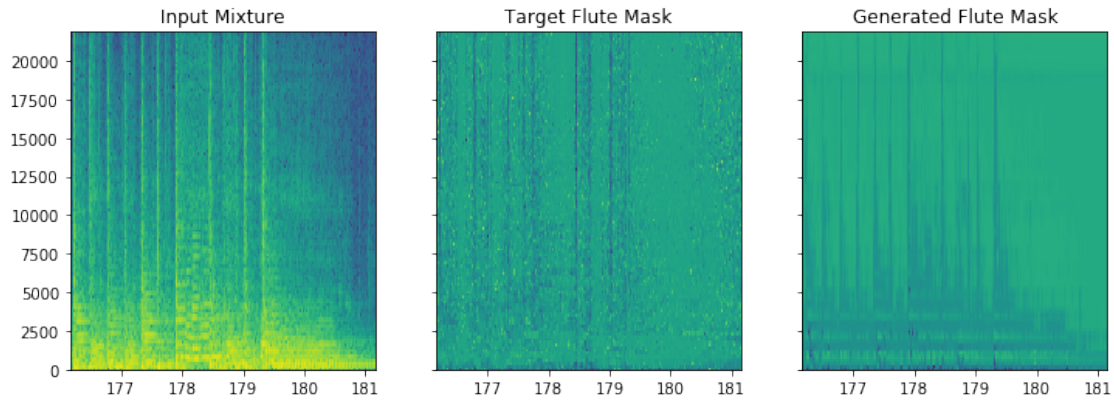
ax1.set_title("Input Mixture")
ax2.set_title("Target Flute Mask")
ax3.set_title("Generated Flute Mask")

```

```
ax1.pcolormesh(times[-num_test:], freqs, 20*np.log10(np.abs(s_mix[:,-num_test:])))
```

```
ax2.pcolormesh(times[-num_test:], freqs, 20*np.log10(np.abs(mask_flute[:,-num_test:])))
ax3.pcolormesh(times[-num_test:], freqs, 20*np.log10(mask_prediction))
```

Out [89]: <matplotlib.collections.QuadMesh at 0x1b489cd3240>



```
In [90]: target = s_mix[:,-num_test:] * mask_flute[:,-num_test:]
```


```
In [91]: prediction = s_mix[:,-num_test:] * mask_prediction
```

```
In [92]: fig, (ax1, ax2, ax3) = plt.subplots(1, 3, sharey=True, figsize=(12, 4))
```

```
ax1.set_title("Input Mixture")
ax2.set_title("Target")
ax3.set_title("Prediction")
```

```
ax1.pcolormesh(times[-num_test:], freqs, 20*np.log10(np.abs(s_mix[:,-num_test:])))
ax2.pcolormesh(times[-num_test:], freqs, 20*np.log10(np.abs(target)))
ax3.pcolormesh(times[-num_test:], freqs, 20*np.log10(np.abs(prediction)))
```

Out [92]: <matplotlib.collections.QuadMesh at 0x1b489ddc668>



4.1 Let's hear it

```
In [99]: _, predicted_audio = signal.istft(prediction, fs=sample_rate)
```

```
In [100]: wavfile.write("fcn_model_flute_prediction.wav", sample_rate, predicted_audio.astype(np.int16))
```

```
In [101]: _, target_audio = signal.istft(target, fs=sample_rate)
```

```
In [102]: wavfile.write("fcn_model_flute_target.wav", sample_rate, target_audio.astype(np.int16))
```

```
In [103]: _, mix_audio = signal.istft(s_mix[:, -num_test:], fs=sample_rate)
```

```
In [104]: wavfile.write("fcn_model_original.wav", sample_rate, mix_audio.astype(np.int16))
```

```
In [ ]:
```

```
In [ ]:
```